

CHIRANJEEVI JOSHI

AI / ML Engineer

+91 78996 59451 · joshichiranjeevi9@gmail.com · [LinkedIn](#) · [GitHub](#) · [Portfolio](#) · healifyai.in

Open to immediate relocation · Bengaluru / international

SUMMARY

AI/ML engineer who took a medical AI product from first commit to live production in 7 months: LLM inference on self-hosted GPUs (vLLM, Qwen3.5-9B, FP4 quantization), an OCR and extraction pipeline for messy Indian lab reports, a clinical evaluation harness with automated regression, and a React Native app, serving external users in 8 Indian languages. Strongest at LLM serving, evaluation harness design, and making model output safe enough for a regulated domain.

EXPERIENCE

Founding Engineer, HealifyAI

Nov 2025 – Present

Self-hosted medical AI platform · healifyai.in · 2-person engineering team

- Took HealifyAI from first commit to a live product in 7 months: explains lab reports in 8 Indian languages for 45+ external users, 230+ reports analyzed (about 5 per user), running entirely on self-hosted dual-GPU infrastructure so no medical data leaves the server.
- Co-built the inference and extraction pipeline: vLLM serving Qwen3.5-9B (FP4-quantized), PaddleOCR as the primary engine with a Qwen-VL fallback on low-confidence pages, and a custom parser for multi-table Indian lab layouts with dynamic column detection and plausibility checks backed by 50+ tests. 98–99% extraction accuracy; responses in 10–15s on the primary path.
- Co-built the clinical validation harness: 5 scoring dimensions over 32 human-verified gold standards spanning 9 lab formats and 6 report types, with automated regression on every change. 97.1% mean end-to-end accuracy; 27 of 32 cases score 100%, the rest root-caused as OCR-ceiling cases on dense tables.
- Owned patient safety: built a 42-keyword histopathology guard with three independent checkpoints after catching the model softening malignant findings into reassurance, plus unit-normalization checks that catch order-of-magnitude errors on critical values.
- Built cross-report intelligence localized to all 8 languages: 40 medical correlation rules, trend detection across reports, and gender-aware filtering, template-based by design so the medical logic cannot hallucinate. Reference ranges prefer lab-printed values, with a 135-entry curated fallback cross-checked against WHO 2024 and RSSDI 2022.
- Co-built the FastAPI backend (SQLAlchemy/Alembic, JWT, PostgreSQL, Neo4j) and a B2B REST API with API-key auth and usage tracking; built the React Native + Expo mobile app (Firebase Auth, EAS builds) and all frontend integration solo.

Full-Stack Developer Intern, Codegnan

Feb 2024 – Mar 2024

Machine Learning Intern, Coincent.ai

Sep 2022 – Nov 2022

PROJECTS

TechDebtAssassin · Python, Claude API, MCP, Streamlit

Feb 2026

- Autonomous code-maintenance agent: scans Python codebases via AST, builds global project context for cross-file type inference, patches missing type hints, and self-verifies with mypy + pytest before re-scanning. CLI, dashboard, and MCP server; CI on every push.

Road Scene Segmentation (IDD) · PyTorch, SegFormer, OpenCV

May 2025

- Fine-tuned SegFormer (MiT-B5) on unstructured Indian road scenes to ~50% mIoU; built the polygon-to-pixel-mask conversion and the full training, inference, and visualization pipeline over 7k+ training and 900+ validation images.

Multilingual NLP Service · Python, T5, Flask

Sep 2025

- Fine-tuned T5 for abstractive summarization and translation, served via Flask REST APIs.

SKILLS

ML / AI: LLM inference & serving (vLLM), LLM evaluation & eval-harness design, LLM agents (tool use, MCP), model quantization (FP4), fine-tuning, OCR pipelines (PaddleOCR, VLM fallback), transformers, PyTorch, Hugging Face

Backend / Infra: Python, FastAPI, SQLAlchemy/Alembic, PostgreSQL, Neo4j, Docker, Docker Compose, Cloudflare, Linux

Frontend / Mobile: React Native, Expo, Next.js, TypeScript, Firebase

EDUCATION

B.E., Artificial Intelligence and Machine Learning

Dec 2020 – Jun 2024

Ballari Institute of Technology and Management · CGPA 7.6